Geometric Limits of Defense Design: Impossibility Theorems for Continuous Prompt Injection Defenses

Manish Bhatt* Sarthak Munshi Vineeth Sai Narajala Idan Habler

Ammar Al-Kahfah Ken Huang Blake Gatto

Abstract

Prior work on LLM safety has focused on discovering adversarial inputs or bounding perturbation robustness around individual examples. We ask a different question: what are the *geometric limits of defense design itself*?

We prove three results of increasing strength. First, any continuous, utility-preserving defense on a connected Hausdorff space must fix boundary points of the alignment deviation surface (*boundary fixation*). Second, when the defense and alignment function are both Lipschitz, the defense cannot uniformly reduce alignment deviation far below the threshold on a positive-measure band around the boundary (*ε-robust constraint*). Third, under a transversality condition—the alignment surface rises faster than the defense can pull it down—the defense leaves a *positive-measure region genuinely unsafe*: $f(D(x)) > \tau$ persists over volume (*persistent unsafe region*).

These are not statements about attacks or training—they are constraints on what any admissible defense mapping can do, independent of the model, data, or optimization procedure. The Tietze extension theorem ensures that any finite set of behavioral observations admits continuous extensions for which all three impossibilities hold.

The paper defines the theoretical speed limit for AI safety. It proves we can't build a 'Perfect Shield' in a continuous space. However, by using multi-layered, discontinuous defenses and ensuring our boundary behavior is benign, we can operate safely within those limits. The 'Manifold of Failure' exists, but we can ensure it's empty of actual harm. The full theory is machine-verified in Lean 4 with Mathlib (35 files, ∼300 theorems, zero `sorry`, three standard axioms) and validated against empirical findings from three frontier LLMs. `https://github.com/mbhatt1/stuff/tree/main/ManifoldProofs`

1 Introduction

Can you build a wrapper around a language model that eliminates all prompt injection vulnerabilities?

Most current work on LLM safety implicitly assumes the answer is yes—at least in principle, even if current tools fall short in practice. The defense mechanisms people have proposed span a wide range: input classifiers that flag suspicious prompts [1], output filters that catch unsafe

*Corresponding author: `manish.bhatt13212@gmail.com`. This work was conducted independently and does not reflect the views of the authors' respective employers.

completions [9], and constitutional rewriting pipelines that reshape model behavior [3]. But despite their differences, all of these defenses share the same underlying structure: they are functions $D\colon X \to X$ that preprocess prompts before the model sees them, aiming to map unsafe inputs to safe equivalents while leaving safe inputs alone.

Prior work has largely focused on discovering or mitigating individual adversarial inputs. We ask a different question: what are the *geometric limits of the defense mapping itself*? We show that under mild and commonly assumed design constraints—continuity and utility preservation—no defense can eliminate the safety boundary. Vulnerabilities arise not from specific inputs, but from the structure of the space and the constraints imposed on admissible defenses. This shifts the problem from adversarial search to geometric inevitability.

We prove three results of increasing strength.

Tier 1: Boundary fixation. If the defense is continuous, utility-preserving, and the prompt space is connected, then the defense must fix at least one boundary point—a prompt where alignment deviation equals the threshold exactly—passing it through with no remediation (Theorem 4.1). The proof is five steps of point-set topology.

Tier 2: ε-robust constraint. If the defense is also K-Lipschitz and the alignment deviation function is L-Lipschitz, the defense cannot uniformly reduce alignment deviation far below τ on a neighborhood of the fixed boundary point. For any x within distance δ of the fixed point z:

$$f(D(x)) \;\geq\; \tau - L(K+1)\,\delta. \tag{1}$$

The ε-band $\{x : \tau - \varepsilon \leq f(x) \leq \tau\}$ has positive measure for every $\varepsilon > 0$ (Theorem 5.2). This constrains the *depth* to which the defense can push near-boundary points into the safe region, though it does not by itself prevent the defense from nudging them slightly below τ.

Tier 3: Persistent unsafe region. Under a transversality condition—the alignment surface has directional slope exceeding $L(K+1)$ at the boundary—the defense leaves a *positive-measure region genuinely unsafe*:

$$f(D(x)) > \tau \quad \text{for all } x \in \mathcal{S}, \qquad \mu(\mathcal{S}) > 0. \tag{2}$$

This is the strongest result: not boundary fixation, not a near-threshold band, but genuine persistent unsafe volume (Theorem 6.2). The condition holds whenever the alignment surface rises faster than the defense can pull it down—empirically observed in all three models we study.

Why continuity? Continuity captures a natural requirement for practical defenses: small changes in input should not produce arbitrarily large changes in the defense's output, as such instability would degrade usability and reliability. Any defense based on learned representations, smooth transformations, or differentiable processing satisfies this constraint. Discontinuous defenses (hard blocklists, keyword filters) escape the theorem but suffer their own well-known failure modes.

From discrete to continuous. All three results apply to continuous interpolants of discrete data. The Tietze extension theorem guarantees that any finite set of behavioral observations in a normal space admits continuous extensions, and the impossibilities hold for every such extension (Theorem 7.1). This does not claim the results hold for the underlying discrete token system directly, but for every continuous model consistent with the observations.

Both the basic and ε-robust proofs are machine-verified in Lean 4 with zero `sorry` statements, as is the persistent unsafe region theorem. All three conditions—continuity, utility preservation, and connectedness—are individually necessary; we give counterexamples for each.

Contributions.

- **Tier 1**: Boundary fixation for continuous utility-preserving defenses (Theorem 4.1).
- **Tier 2**: ε-robust constraint showing the defense cannot uniformly reduce AD far below τ on a positive-measure band (Theorems 5.1 and 5.2).
- **Tier 3**: Persistent unsafe region under transversality (Theorem 6.2).
- A formally verified bridge from discrete data to continuous theory via Tietze extension (Theorem 7.1).
- Counterexamples proving each hypothesis is necessary (Section 4.1).
- Supporting theorems: basin structure, fragmentation, convergence, transferability, authority monotonicity, gradient ascent, fine-tuning stability, cost asymmetry (Sections 8 to 11).
- Theory–experiment correspondence with three LLMs (Section 12).
- Lean 4 artifact: 35 files, ∼300 theorems, zero `sorry`, three standard axioms (Section 14).

2 Related Work

Adversarial robustness. The existence of adversarial examples in neural networks has been known for over a decade [15, 8], and a substantial body of work has pursued formal verification of robustness properties—through abstract interpretation [14], SMT-based methods [10], and randomized smoothing [7]. Unlike adversarial robustness results, which characterize the existence of perturbations or bound robustness radii around individual inputs under specific norms, our result is a *structural constraint on the defense mapping itself*, independent of perturbation models, norms, or training procedures. Certified defenses provide guarantees about perturbation neighborhoods; our result constrains the *global behavior* of the defense function, showing that no such defense can eliminate boundary regions.

Prompt injection and jailbreaking. Techniques for finding individual vulnerabilities include gradient-based attacks like GCG [18], search-based approaches like PAIR [6] and TAP [11], and quality-diversity methods like Rainbow Teaming [13]. The Manifold of Failure framework [4] uses MAP-Elites to map the *continuous topology* of failures. Existing work identifies individual failure cases or distributions of attacks. Our result shows that these failures are not incidental but *topologically unavoidable* under mild assumptions, providing a formal explanation for their persistence.

Impossibility results in ML. The no-free-lunch theorems [17] limit average performance across tasks. The robustness–accuracy tradeoff literature [16] shows that adversarial robustness costs standard performance. Our result is different in character: it holds *pointwise* for any continuous defense on a *fixed task*, with no dependence on distributions, learning algorithms, or training data.

Topological views of neural networks. Work on decision boundary geometry, the manifold hypothesis, and topology of neural representations has established that topological structure matters. While our proof relies on elementary topology, the novelty lies in identifying that *standard design constraints for LLM defenses induce a forced fixed-point structure at the safety boundary*—a connection not previously characterized. The persistent unsafe region theorem (Tier 3) goes further, using metric structure to establish volumetric persistence.

Formal verification of ML systems. Proof assistants like Lean and Coq have been applied to verify properties of neural networks [2] and optimization algorithms [5]. To the best of our knowledge, ours is the first formally verified impossibility result specifically targeting prompt injection defense.

3 Setup

Definition 3.1 (Alignment Deviation Function). *Let X be a topological space. An* alignment deviation function *is a continuous map $f\colon X \to \mathbb{R}$. Given a threshold $\tau \in \mathbb{R}$:*

$$S_\tau = \{x \in X : f(x) < \tau\} \qquad \textit{(safe region)} \tag{3}$$

$$U_\tau = \{x \in X : f(x) > \tau\} \qquad \textit{(unsafe region)} \tag{4}$$

$$B_\tau = \{x \in X : f(x) = \tau\} \qquad \textit{(boundary)} \tag{5}$$

Definition 3.2 (Defense). *A* defense *is a continuous map $D\colon X \to X$. It is* utility-preserving *if $D(x) = x$ for all $x \in S_\tau$, and* complete *if $f(D(x)) < \tau$ for all $x \in X$.*

4 Tier 1: Boundary Fixation

Theorem 4.1 (Boundary Fixation). *Let X be a connected Hausdorff space. Let $f\colon X \to \mathbb{R}$ be continuous with $S_\tau \neq \emptyset$ and $U_\tau \neq \emptyset$. Let $D\colon X \to X$ be continuous with $D|_{S_\tau} = \mathrm{id}$.*

Then there exists $z \in X$ with $f(z) = \tau$ and $D(z) = z$. Moreover, every $z \in \overline{S_\tau} \setminus S_\tau$ satisfies $f(z) = \tau$ and $D(z) = z$. This set is nonempty.

Proof. Five steps, each using exactly one hypothesis.

Step 1 (Fixed-point set is closed). $\mathrm{Fix}(D) = \{x : D(x) = x\}$ is the preimage of the diagonal $\Delta \subset X \times X$ under $x \mapsto (D(x), x)$. Hausdorff $\Rightarrow \Delta$ closed $\Rightarrow \mathrm{Fix}(D)$ closed.

Step 2 (Closure of safe region $\subseteq$ fixed points). $S_\tau \subseteq \mathrm{Fix}(D)$ by utility preservation. $\mathrm{Fix}(D)$ closed $\Rightarrow \overline{S_\tau} \subseteq \mathrm{Fix}(D)$.

Step 3 (Safe region is not closed). $S_\tau = f^{-1}((-\infty, \tau))$ is open. If also closed, it would be clopen, hence $\emptyset$ or X—contradicting both regions nonempty.

Step 4 (Boundary point in closure). S_τ not closed $\Rightarrow \overline{S_\tau} \supsetneq S_\tau$. Any $z \in \overline{S_\tau} \setminus S_\tau$ satisfies $f(z) \leq \tau$ (closure $\subseteq \{f \leq \tau\}$) and $f(z) \geq \tau$ ($z \notin S_\tau$). So $f(z) = \tau$.

Step 5 (Defense fixes it). $z \in \overline{S_\tau} \subseteq \mathrm{Fix}(D)$, so $D(z) = z$ and $f(D(z)) = \tau \not< \tau$. □

Corollary 4.2. *No continuous utility-preserving defense is complete.*

4.1 Each Hypothesis Is Necessary

Counterexample 4.3 (Removing connectedness). *$X = \{0, 1\}$ discrete, $f(0) = 0$, $f(1) = 1$, $\tau = 0.5$. $D(0) = 0$, $D(1) = 0$: continuous, utility-preserving, complete.*

Counterexample 4.4 (Removing continuity). *$X = [0, 1]$, $f(x) = x$, $\tau = 0.5$. $D(x) = x$ for $x < 0.5$, $D(x) = 0$ for $x \geq 0.5$: utility-preserving and complete, but discontinuous at 0.5.*

Counterexample 4.5 (Removing utility preservation). *$D(x) = x_0$ for a fixed safe point: continuous and complete, but destroys all inputs.*

Remark 4.6 (The defense trilemma). *A defense can satisfy at most two of: continuity, utility preservation, completeness.*

Remark 4.7 (Not all boundary points are fixed). *The theorem captures $\overline{S_\tau} \setminus S_\tau$, not all of B_τ. Boundary points isolated from the safe side may escape. Verified in Lean with $f(x) = x^2$, $\tau = 0$.*

5 Tier 2: ε-Robust Constraint

Tier 1 fixes boundary points—potentially measure zero. Tier 2 extends the constraint to a positive-measure neighborhood.

Theorem 5.1 (ε-Robust Defense Constraint). *Let X be a connected Hausdorff metric space, f be L-Lipschitz with $S_\tau, U_\tau \neq \emptyset$, and D be K-Lipschitz with $D|_{S_\tau} = \mathrm{id}$.*

Then there exists $z \in X$ with $f(z) = \tau$, $D(z) = z$, such that for all $x \in X$ (with strongest constraint near z):

$$f(D(x)) \;\geq\; \tau \;-\; L(K+1)\,\mathrm{dist}(x,z). \tag{6}$$

Proof. From Theorem 4.1, z exists. Since D is K-Lipschitz and $D(z) = z$: $\mathrm{dist}(D(x), x) \leq (K+1)\,\mathrm{dist}(x,z)$ (triangle inequality). Since f is L-Lipschitz: $|f(D(x)) - f(x)| \leq L(K+1)\,\mathrm{dist}(x,z)$. Extracting the lower bound gives $f(D(x)) \geq f(x) - L(K+1)\,\mathrm{dist}(x,z) \geq \tau - L(K+1)\,\mathrm{dist}(x,z)$ (using $f(z) = \tau$ and Lipschitz of f). □

Theorem 5.2 (Positive-Measure ε-Band). *Under the hypotheses of Theorem 5.1, for every $\varepsilon > 0$ such that f takes values below $\tau - \varepsilon$, the ε-band $\mathcal{B}_\varepsilon = \{x : \tau - \varepsilon \leq f(x) \leq \tau\}$ has positive measure (under any measure positive on nonempty open sets). Specifically,* $\mathrm{ball}(c, \varepsilon/(4L)) \subseteq \mathcal{B}_\varepsilon$ *for any c with $f(c) = \tau - \varepsilon/2$.*

Remark 5.3 (What Tier 2 does and does not say). *Theorem 5.1 constrains the* depth *to which the defense can push near-boundary points below τ. It does* not *prove these points remain unsafe—the defense may map them to values in $[\tau - L(K+1)\delta,\ \tau)$, which is below threshold. The bound is tight: a defense mapping every point in a δ-ball to f-value $\tau - L(K+1)\delta$ would saturate it.*

6 Tier 3: Persistent Unsafe Region

Tier 2 constrains depth but permits the defense to nudge points below τ. Under a transversality condition, Tier 3 proves the defense *cannot make certain points safe at all.*

Definition 6.1 (Steep region). *Given a fixed boundary point z (from Theorem 4.1), define*

$$\mathcal{S} = \{x \in X : f(x) > \tau + L(K+1)\ \mathrm{dist}(x,z)\}.$$

This is the set of points where the alignment deviation exceeds τ by more than the defense's Lipschitz bound can compensate.

Theorem 6.2 (Persistent Unsafe Region). *Under the hypotheses of Theorem 5.1, if $\mathcal{S}$ is nonempty, then:*

1. *$\mathcal{S}$ is open.*
2. *$\mathcal{S}$ has positive measure (under any measure positive on nonempty open sets).*
3. *For every $x \in \mathcal{S}$: $f(D(x)) > \tau$.*

In particular, $\mu(\{x : f(D(x)) > \tau\}) > 0$—the defense leaves a positive-measure region genuinely unsafe.

Proof. **(1)** $\mathcal{S}$ is the strict superlevel set of the continuous function $x \mapsto f(x) - L(K+1)\,\mathrm{dist}(x,z)$ at level τ, hence open.

(2) Open and nonempty implies positive measure.

(3) For $x \in \mathcal{S}$: $f(x) > \tau + L(K+1)\,\mathrm{dist}(x,z)$. From the proof of Theorem 5.1, the defense satisfies $f(D(x)) \geq f(x) - L(K+1)\,\mathrm{dist}(x,z)$. Combining:

$$f(D(x)) \geq f(x) - L(K+1)\,\mathrm{dist}(x,z) > \tau.$$

The defense cannot bring these points below threshold. □

Remark 6.3 (When is $\mathcal{S}$ nonempty?)**.** *The set $\mathcal{S}$ is nonempty whenever the alignment surface rises with directional slope exceeding $L(K+1)$ at the boundary point z. In a normed space, if f has directional derivative $c > L(K+1)$ along some unit vector v at z, then $z + tv \in \mathcal{S}$ for small $t > 0$ (verified in Lean as* `transversality_from_deriv`*).*

Informally: if the alignment surface rises faster than the defense can pull it down, unsafe behavior persists over volume. This condition is empirically observed in all three models we study—the Lipschitz constant of the alignment surface near boundary crossings consistently exceeds reasonable defense aggressiveness bounds.

Remark 6.4 (The defense's fundamental dilemma)**.** *The three tiers reveal a progressive bind. Tier 1: the defense* must *fix boundary points. Tier 2: it cannot push nearby points* deep *into safety. Tier 3: if the landscape is steep, it cannot push them into safety* at all. *Increasing K (defense aggressiveness) raises the transversality threshold $L(K+1)$, requiring an even steeper landscape to trigger Tier 3—but it also makes the defense less stable and widens the Tier 2 constraint band. The defense cannot escape the geometry.*

7 From Discrete Data to Continuous Theory

Theorem 7.1 (Continuous Relaxation)**.** *Let X be a connected, normal, Hausdorff space. Let $S \subset X$ be a finite set of behavioral observations with $g(p) < \tau$ for some p and $g(q) > \tau$ for some q.*

Then there exists a continuous function $f\colon X \to \mathbb{R}$ agreeing with all observations for which Theorems 4.1 and 5.1 hold. If f is additionally Lipschitz (as it is for GP posterior means under standard kernel regularity assumptions, e.g., jointly continuous kernel on a compact domain), Theorem 6.2 applies wherever the transversality condition is met.

Proof. Finite subsets of T_1 spaces are closed. By the Tietze extension theorem, g extends to a continuous f agreeing at every sample point. Since $f(p) < \tau < f(q)$, Theorem 4.1 applies. Lipschitz extensions (which exist by McShane–Whitney) enable Theorems 5.1 and 6.2. □

Remark 7.2. *The interpolant is not unique. All three impossibilities hold for* every *continuous extension consistent with the data. This does not claim the results hold for the underlying discrete token system directly, but for every continuous model consistent with the observations.*

Corollary 7.3. *Theorem 7.1 applies in $\mathbb{R}^n$ for every $n \geq 1$.*

8 The Vulnerability Landscape

Theorem 8.1 (Basin Structure)**.** *If f is continuous and $f(p) > \tau$, then U_τ is open. Under any measure that assigns positive mass to nonempty open sets (e.g., Lebesgue measure on $\mathbb{R}^n$), U_τ has positive measure.*

Theorem 8.2 (Basin Fragment Minimum Size). *If f is L-Lipschitz with $f(p) > \tau$, the connected component of U_τ containing p has diameter $\geq 2(f(p) - \tau)/L$.*

9 Attack Properties

Theorem 9.1 (Perturbation Robustness). *If f is L-Lipschitz and $f(p) > \tau$, then* $\mathrm{ball}(p, (f(p) - \tau)/L) \subseteq U_\tau$. *The radius is monotone in $f(p)$.*

Theorem 9.2 (Iterative Convergence). *Any monotone-improvement operator T with score in $[0, 1]$ converges. If each step gains $\geq \delta$, convergence takes $\leq \lfloor 1/\delta \rfloor$ steps.*

Theorem 9.3 (Transferability). $\|f - g\|_\infty \leq \delta \implies \{f > \tau + \delta\} \subseteq \{g > \tau\}$. *Transfer costs zero additional queries.*

Theorem 9.4 (Authority Monotonicity). *If $f(y, \cdot)$ is monotone non-decreasing in authority for each fixed indirection y, then the vulnerability set $\{x : f(y, x) > \tau\}$ is upward-closed and a critical threshold $a_2^*(y)$ exists (by IVT, when $f(y, \cdot)$ is continuous and crosses τ). If additionally f is monotone non-decreasing in indirection for each fixed authority, the threshold curve $y \mapsto a_2^*(y)$ is non-increasing.*

Theorem 9.5 (Gradient Ascent). *If f has nonzero Fréchet derivative at x, there exist v and $\varepsilon > 0$ with $f(x + \varepsilon v) > f(x)$.*

10 Stability Under Fine-Tuning

Theorem 10.1 (Interior Stability). *If $\|f - g\|_\infty \leq \varepsilon$: $f(x) > \tau + \varepsilon \implies g(x) > \tau$, and $f(x) < \tau - \varepsilon \implies g(x) < \tau$. Only the band $|f(x) - \tau| \leq \varepsilon$ is uncertain.*

Theorem 10.2 (Crossing Preservation). *If f crosses τ transversally on $[a, b]$ and g is continuous with $\|f - g\|_\infty \leq \varepsilon$, then g also crosses τ on $[a, b]$.*

Theorem 10.3 (Patching Is Nonlocal). *There exist f, g with $\|f - g\|_\infty \leq \varepsilon$ such that eliminating a vulnerability at one point necessarily changes values at distant points. (Verified by explicit construction in Lean.)*

11 Cost Asymmetry

Theorem 11.1 (Exponential Cost Asymmetry for Grid-Based Defense). *For a defense monitoring the behavioral space via a grid with $N \geq 2$ bins per axis in d dimensions: attack cost $\leq 1/\delta$ (dimension-independent); grid-based defense cost $= N^d$ (exponential in d); the ratio $\delta \cdot N^d \to \infty$ as $d \to \infty$.*

At $d = 2$ with $N = 25$ and $\delta = 0.01$, the ratio is 6.25. At $d = 10$, it climbs to $\sim 10^{12}$. Verified by `norm_num` in Lean.

12 Experimental Validation

We validate against the Manifold of Failure framework [4], which maps three LLMs over a 2D behavioral space (query indirection $\times$ authority framing).

Table 1: Falsifiable predictions confirmed by empirical data.

Theorem	Predicts	Confirmed by
Basin Structure (8.1)	Basins are open with positive measure	Heatmaps show extended regions
Fragmentation (8.2)	Smooth → large basins; rough → mosaic	Llama: mesa; GPT-OSS: mosaic
Convergence (9.2)	Attacks exhibit monotone convergence	Convergence curves plateau
Transferability (9.3)	Similar surfaces → shared basins	Llama .93 → GPT-OSS .73 → Mini .47
Authority (9.4)	Horizontal banding	Bands at $a_2 \approx$ .25–.35, .65–.85
Persistent (6.2)	Steep boundaries → unsafe volume persists	Llama's .93 plateau immune to defense
Interior Stability (10.1)	Deep basin points survive fine-tuning	Vulnerability persists across variants
Cost (11.1)	2D tractable, high-d intractable	15K queries fill 63% at $d=2$

Llama-3-8B (mean AD 0.93, basin rate 93.9%): near-flat mesa, small L, large robustness radii. The directional slope of the alignment surface at boundary crossings far exceeds $L(K+1)$ for any reasonable K, placing this model squarely in Tier 3 territory: the defense provably cannot eliminate the vulnerability volume.

GPT-OSS-20B (mean AD 0.73, basin rate 64.3%): rugged landscape, large L, small fragments. Horizontal bands confirm authority monotonicity. The fragmented boundary presents many transversal crossings, each generating a persistent unsafe pocket.

GPT-5-Mini (peak AD 0.50, basin rate 0%): at $\tau = 0.50$, $U_\tau = \emptyset$, so none of the three tiers apply—correctly predicting no impossibility. The model's strategy: make the boundary harmless rather than eliminate it.

Quantitative checks. Robustness radius: $L \approx 5$ for Llama gives $r = (0.93-0.50)/5 = 0.086$, covering $\sim$2 grid cells—consistent with heatmap smoothness. Cost ratio: $0.01 \times 625 = 6.25$ theoretical vs. $\sim$4 empirical.

13 The Engineering Prescription

Make the boundary shallow. Set τ so that boundary-level behavior is benign (GPT-5-Mini: AD $= 0.50$).

Reduce the Lipschitz constant. Smaller L makes the surface predictable and raises the transversality threshold for Tier 3.

Reduce the effective dimension. Defense cost grows as N^d; constraining the prompt interface reduces d.

Monitor the boundary. Theorem 10.2 shows transversal crossings persist under fine-tuning. Rather than eliminate them, monitor approach to the boundary.

14 The Lean Artifact

The complete theory is verified in Lean 4.28.0 with Mathlib v4.28.0. The artifact comprises 35 files:

- 10 core theory files (MoF_01–MoF_10)
- 10 cost theory files (MoF_Cost_01–MoF_Cost_10)
- 10 advanced theory files (MoF_Adv_01–MoF_Adv_10)
- 1 continuous relaxation (MoF_ContinuousRelaxation)
- 1 ε-robust + persistent unsafe region (MoF_11_EpsilonRobust)
- 3 capstone files (MasterTheorem, Euclidean instantiation, verification)

Every theorem depends on exactly three axioms: `propext`, `Classical.choice`, `Quot.sound`—Lean's standard foundations. No additional axioms.

Key Mathlib lemmas used:

Lean name	**Role**
`isClosed_diagonal`	Hausdorff $\Rightarrow$ diagonal closed
`LipschitzWith.dist_le_mul`	Lipschitz distance bound
`isClopen_iff`	Connected $\Rightarrow$ no proper clopen
`closure_minimal`	Closure in closed superset
`IsOpen.measure_pos`	Open nonempty $\Rightarrow$ positive measure
`intermediate_value`$_2$	IVT on connected spaces

15 Limitations

Tier 1 is at the boundary. Fixed points satisfy $f(z) = \tau$ exactly. If τ-level behavior is benign, the theorem is true but harmless.

Tier 2 constrains depth, not direction. The defense may push near-boundary points slightly below τ. The bound limits how far, not whether.

Tier 3 requires transversality. The persistent unsafe region exists only where the alignment surface is steep ($c > L(K + 1)$). Flat or gently sloping boundaries escape Tier 3 (though not Tiers 1–2).

The Tietze bridge applies to interpolants. The impossibilities hold for every continuous extension of discrete data, not for the underlying discrete system itself.

Grid-based cost asymmetry. Theorem 11.1 assumes exhaustive grid enumeration. Learning-based defenses that generalize across the space may sidestep the exponential bound.

Single-turn only. The manifold may shift across conversation turns.

16 Conclusion

We have established three impossibility results for prompt injection defense, forming a progressive hierarchy. Tier 1 (boundary fixation): any continuous, utility-preserving defense must fix boundary points. Tier 2 (ε-robust constraint): the defense cannot uniformly reduce alignment deviation far below τ on a positive-measure band. Tier 3 (persistent unsafe region): under transversality, the defense leaves positive-measure volume genuinely unsafe—$f(D(x)) > \tau$ persists.

The key insight is not about attacks or training but about the *geometry of the defense itself*: continuity and utility preservation together create an anchor at the boundary that constrains what the defense can do in its neighborhood, and metric structure (Lipschitz bounds) makes this constraint quantitative and volumetric.

The discrete-continuous gap is addressed via Tietze extension: every continuous interpolant of discrete data inherits all three impossibilities. Supporting theorems show that vulnerabilities come in basins, attacks converge and transfer, vulnerabilities survive fine-tuning, and defense cost grows exponentially. Each result is validated against empirical data from three frontier LLMs.

The practical prescription: make the boundary shallow (raise τ), smooth (reduce L), and low-dimensional (constrain the interface). GPT-5-Mini's ceiling at AD = 0.50 shows this working in practice.

The full theory is machine-verified in Lean 4: 35 files, approximately 300 theorems, zero `sorry`, three standard axioms.

References

[1] G. Alon and M. Kamfonas. Detecting language model attacks with perplexity. *arXiv preprint arXiv:2308.14132*, 2023.

[2] A. Bagnall and G. Stewart. Certifying the true error: Machine learning in Coq with verified generalization guarantees. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2019.

[3] Y. Bai et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022.

[4] M. Bhatt, S. Munshi, V. S. Narajala, I. Habler, A. Al-Kahfah, K. Huang, and B. Gatto. Manifold of failure: Behavioral attraction basins in language models. *arXiv preprint arXiv:2602.22291v2*, 2026.

[5] J. Bolte et al. Formal verification of optimization algorithms. *Mathematical Programming*, 2024.

[6] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong. Jailbreaking black box large language models in twenty queries. *arXiv preprint arXiv:2310.08419*, 2024.

[7] J. Cohen, E. Rosenfeld, and J. Z. Kolter. Certified adversarial robustness via randomized smoothing. *Proceedings of ICML*, 2019.

[8] I. J. Goodfellow, J. Shlens, and C. Szegedy. Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*, 2014.

[9] H. Inan et al. Llama Guard: LLM-based input-output safeguard for human-AI conversations. *arXiv preprint arXiv:2312.06674*, 2023.

[10] G. Katz, C. Barrett, D. L. Dill, K. Julian, and M. J. Kochenderfer. Reluplex: An efficient SMT solver for verifying deep neural networks. *Proceedings of CAV*, 2017.

[11] A. Mehrotra et al. Tree of attacks: Jailbreaking black-box LLMs automatically. *arXiv preprint arXiv:2312.02119*, 2024.

[12] J.-B. Mouret and J. Clune. Illuminating search spaces by mapping elites. *arXiv preprint arXiv:1504.04909*, 2015.

[13] M. Samvelyan et al. Rainbow teaming: Open-ended generation of diverse adversarial prompts. *Advances in Neural Information Processing Systems*, 37, 2024.

[14] G. Singh, T. Gehr, M. Püschel, and M. Vechev. An abstract domain for certifying neural networks. *Proceedings of POPL*, 2019.

[15] C. Szegedy et al. Intriguing properties of neural networks. *arXiv preprint arXiv:1312.6199*, 2013.

[16] D. Tsipras, S. Santurkar, L. Engstrom, A. Turner, and A. Madry. Robustness may be at odds with accuracy. *arXiv preprint arXiv:1805.12152*, 2018.

[17] D. H. Wolpert and W. G. Macready. No free lunch theorems for optimization. *IEEE Trans. Evol. Comput.*, 1(1):67–82, 1997.

[18] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.